Tanmay Singh

Delhi, India | tanmay.work10@gmail.com | LinkedIn | Github | Portfolio

Data Engineer with experience designing and scaling ETL pipelines, relational schemas, and cloud data infrastructure for AI-driven products and enterprise financial systems. Skilled in building reliable, high-throughput data platforms across Airflow, dbt, BigQuery, PostgreSQL, Supabase, and AWS, with a track record of cutting pipeline runtime by 40% and processing 10M+ records monthly at production scale.

EDUCATION

Arizona State University — Master of Science in Data Science, Analytics and Engineering
Manipal University, Jaipur — Bachelor of Technology in Computer Science and Engineering

EXPERIENCE

Vault IQ **Dec 2025 – May 2026**
Data Engineer (AI and Backend Systems) *Remote*

- Designed and implemented relational schemas and storage for user, brand, product, and social data in Supabase (PostgreSQL), giving 4+ product modules one structured data flow
- Worked full stack across the data layer, backend services, and product UI of an AI-powered brand development platform integrating OpenAI, Gemini, and Flux Pro
- Sped up the brand-generation pipeline that orchestrates multi-model LLM and image-generation calls for content and product visualization
- Built a photoshoot agent that stages generated product images in marketing-ready settings such as hero shots, using reasoning-based prompt orchestration while preserving product fidelity
- Built and deployed a production chatbot with a custom knowledge-base integration layer for contextual querying and FAQ automation
- Built the membership dashboard backend serving brand portfolio, product catalog, and revenue metrics in real time, integrating 6+ external services and APIs

Capgemini Pvt Ltd **Oct 2020 – Mar 2023**
Data Engineer – HSBC UK (Application Data Management) *United Kingdom*

- Designed and orchestrated scalable ETL pipelines in Apache Airflow (DAG-based workflows) for batch processing of HSBC's customer credit data systems
- Managed and processed 10M+ customer records monthly across Google BigQuery and PostgreSQL environments, ensuring data consistency, integrity, and availability
- Developed transformation and validation workflows in Python and PySpark, with data quality checks aligned to business-critical credit risk models
- Optimized BigQuery queries and storage structures and implemented SCD logic, reducing batch pipeline runtime by 40% and meeting ETL SLAs
- Built monitoring into pipeline runs to catch data anomalies pre-emptively, improving the reliability of production credit data systems

PROJECTS

CheckYourAQI | Air Quality Forecasting Platform (live)

- Built an hourly pipeline on AWS: Airflow on EC2 lands air-quality, weather-forecast, and satellite fire data in an S3/Iceberg lakehouse, modeled by 27 dbt models on Athena with 44 data tests
- Serves a public dashboard and FastAPI on Lambda from versioned Parquet snapshots; weekly retraining on GitHub Actions; runs unattended for ~$21/month with failure alerts

Mettle | Gamified Workout Tracker (in development)

- Offline-first React Native app: local SQLite synced to Supabase Postgres with row-level security; a shared TypeScript engine recomputes XP server-side so leaderboards can't be edited on the phone

Real-Time Clickstream Analytics Pipeline on AWS

- Serverless streaming pipeline: Kinesis and Lambda ingest events to S3, Glue ETL partitions by date, Athena queries the result, and Step Functions orchestrates the run with SNS alerts

SKILLS

Languages: Python, SQL, PySpark, TypeScript
Data Platforms & Warehousing: PostgreSQL, Google BigQuery, Supabase, Amazon S3, Apache Iceberg, Athena, Data Modeling, Schema Design
Pipelines & Orchestration: Apache Airflow (DAG-based ETL), dbt, AWS Glue, Kinesis, Step Functions, Batch & Streaming, Data Quality Checks
Cloud & DevOps: AWS (EC2, Lambda), Google Cloud, Docker, GitHub Actions, Git
Other: FastAPI, REST API Integration, LLM/Generative AI Integration, Cross-functional Collaboration